AI Evals for Engineers & PMs

Know where your AI fails.
Learn how to improve it.

4 weeks / Live online / English

21 November - 18 December 2026 / **$500 USD**

A repeatable way to improve AI products

Build an agent, inspect its traces, learn from expert reviews, and turn recurring failures into evaluations. Use those evaluations to choose improvements and catch regressions.

What you will practice

- Inspect messages, retrieval and tool calls in one trace.
- Create a reviewed dataset with cases that challenge your product.
- Test code checks and LLM judges against domain-expert labels.
- Compare quality, cost and latency before shipping a change.

With Ion Mosnoi

11+ years in AI/ML. Production AI work with startups and enterprises; research in multi-agent coordination and computer vision in France and Finland.

For engineers, PMs and domain experts. Engineers should be comfortable with basic Python; PMs can pair for coding labs. Plan around 3-5 hours per week. Exact live-session times will be confirmed by email; contact Ion before payment if timing affects your attendance.

WEEK 01 / 21 - 27 NOVEMBER

Build an agent you can inspect

- Define the task, tool contracts, allowed actions and escalation rules.
- Build a small agent loop with tool calls, stop conditions and bounded retries.
- Log messages, retrieval, tool parameters, outputs, cost and latency.
- Build a trace viewer that makes the failure easy to inspect.

Lab outcome Run a support agent. Open one trace, locate a mistake and point to the evidence.

WEEK 02 / 28 NOVEMBER - 4 DECEMBER

Turn reviews into a test dataset

- Choose a domain expert. Annotate Good/Bad with a specific reason and evidence.
- Review across roles and conditions; group recurring mistakes into an error taxonomy.
- Create cases from features x roles x scenarios, including clarification and refusal.
- Check synthetic labels, preserve sources and separate train, dev and test cases.

Lab outcome Review a batch of traces, define your error types and build a versioned dataset.

WEEK 03 / 5 - 11 DECEMBER

Build evaluators you can trust

- Write one pass/fail evaluator per failure mode; use code for objective checks.
- Give a judge a clear rule, contrasting examples and an evidence-based reason.
- Compare with expert labels: inspect false positives, false negatives, precision and recall.
- Improve the judge on dev cases. Keep test cases untouched and report results by category.

Lab outcome Calibrate a judge against human reviews. Explain the disagreements instead of hiding them in an average.

WEEK 04 / 12 - 18 DECEMBER

Improve the product. Catch regressions.

- Diagnose failures by data, retrieval, prompt, model or tool flow.
- Try focused fixes: better chunks, query rewriting, examples, routing or smaller task scopes.
- Compare against a baseline on the same cases, including quality, cost and latency.
- Run evals in CI, set release checks and bring production failures back into the dataset.

Lab outcome Present a before/after experiment: what changed, which errors improved and what still fails.