AI Evals for Engineers & PMs

How do you know your AI actually works?

Learn to find failure patterns, build evaluations you can trust, and improve accuracy, reliability and cost.

4 weeks / Live online / English

21 November - 18 December 2026 / **$500 USD**

Included in your enrollment

- Live sessions and practical labs.
- 10+ hours of live office hours for questions and feedback.
- Private Discord access for discussion between sessions.
- Lifetime access to course recordings.
- Prompts, datasets and evaluation code to keep using.

With Ion Mosnoi

11+ years in AI/ML. Production AI work with startups and enterprises; research in multi-agent coordination and computer vision in France and Finland.

For engineers, PMs and domain experts. Basic Python for engineers; PMs can pair for coding labs. Plan around 3-5 hours weekly for core sessions and practice, plus optional office hours.

Exact live-session and office-hour times will be confirmed by email. Contact Ion before payment if timing affects your attendance.

WEEK 01 / 21 - 27 NOVEMBER

Build an agent you can inspect

- Define the task, tool contracts, allowed actions and escalation rules.
- Build a small agent loop with tool calls, stop conditions and bounded retries.
- Log messages, retrieval, tool parameters, outputs, cost and latency.
- Build a trace viewer that makes the failure easy to inspect.

Lab outcome Inspect an agent run, identify a failure and point to the evidence.

WEEK 02 / 28 NOVEMBER - 4 DECEMBER

Turn reviews into a test dataset

- Choose a domain expert. Annotate Good/Bad with a specific reason and evidence.
- Review across roles and conditions; group recurring mistakes into an error taxonomy.
- Create cases from features x roles x scenarios, including clarification and refusal.
- Check synthetic labels, preserve sources and separate train, dev and test cases.

Lab outcome Review a batch of traces, define your error types and build a versioned dataset.

WEEK 03 / 5 - 11 DECEMBER

Build evaluators you can trust

- Write one pass/fail evaluator per failure mode; use code for objective checks.
- Give a judge a clear rule, contrasting examples and an evidence-based reason.
- Compare with expert labels: inspect false positives, false negatives, precision and recall.
- Improve the judge on dev cases. Keep test cases untouched and report results by category.

Lab outcome Calibrate a judge against human reviews. Explain the disagreements instead of hiding them in an average.

WEEK 04 / 12 - 18 DECEMBER

Improve the product. Catch regressions.

- Diagnose failures by data, retrieval, prompt, model or tool flow.
- Try focused fixes: better chunks, query rewriting, examples, routing or smaller task scopes.
- Compare against a baseline on the same cases, including quality, cost and latency.
- Run evals in CI, set release checks and bring production failures back into the dataset.

Lab outcome Present a before/after experiment: what changed, which errors improved and what still fails.