AI Evals for Engineers & PMs

How do you know your AI actually works?

Find failure patterns, build evaluations you can trust, and improve accuracy, reliability and cost.

November: 21 Nov - 18 Dec 2026 / **EUR 500**

December: 19 Dec 2026 - 15 Jan 2027 / **EUR 1,000**

4 weeks each / Live online / English

Included in your enrollment

- Live sessions and 10+ hours of live office hours.
- Lifetime access to course recordings, materials and private Discord.
- Future live editions of this course at no extra course fee.
- 200+ page course reader, prompts, datasets and evaluation code.
- Six assignments with worked solutions and recorded walkthroughs.
- Completion certificate after instructor review of required coursework.

With Ion Mosnoi

11+ years in AI/ML. Production AI work with startups and enterprises; research in multi-agent coordination and computer vision in France and Finland.

Plan around 3-5 hours weekly for core sessions and practice, plus optional office hours. Exact live times will be confirmed by email; check with Ion before paying if needed.

WEEK 01

Build an agent you can inspect

- Define the task, tool contracts, allowed actions and escalation rules.
- Build a small agent loop with tool calls, stop conditions and bounded retries.
- Log messages, retrieval, tool parameters, outputs, cost and latency.
- Build a trace viewer that makes the failure easy to inspect.

Lab outcome: Inspect an agent run, identify a failure and point to the evidence.

WEEK 02

Turn reviews into a test dataset

- Choose a domain expert. Annotate Good/Bad with a specific reason and evidence.
- Review across roles and conditions; group recurring mistakes into an error taxonomy.
- Create cases from features x roles x scenarios, including clarification and refusal.
- Check synthetic labels, preserve sources and separate train, dev and test cases.

Lab outcome: Review a batch of traces, define your error types and build a versioned dataset.

WEEK 03

Build evaluators you can trust

- Write one pass/fail evaluator per failure mode; use code for objective checks.
- Give a judge a clear rule, contrasting examples and an evidence-based reason.
- Compare with expert labels: inspect false positives, false negatives, precision and recall.
- Improve the judge on dev cases. Keep test cases untouched and report results by category.

Lab outcome: Calibrate a judge against human reviews. Explain the disagreements instead of hiding them in an average.

WEEK 04

Improve the product. Catch regressions.

- Diagnose failures by data, retrieval, prompt, model or tool flow.
- Try focused fixes: better chunks, query rewriting, examples, routing or smaller task scopes.
- Compare against a baseline on the same cases, including quality, cost and latency.
- Red-team prompt injection, unsafe tool calls and missing permissions; add human approval where needed.
- Run evals in CI. Monitor quality, drift, cost and latency, then bring production failures back into the dataset.

Lab outcome: Present a before/after experiment: what changed, which errors improved and what still fails.

Practice, support and next steps.

Six assignments

- Inspect an agent trace and make its behavior visible.
- Review traces, write expert notes and group errors.
- Create a synthetic dataset with reviewed labels and safe splits.
- Write one binary evaluator per error.
- Calibrate a judge against domain-expert labels.
- Run a controlled improvement experiment and define release checks.

Who this is for

Engineers, product managers, domain experts and technical leaders. Basic Python is useful for coding labs; PMs can pair. Use an authorized, anonymized work project or the supplied practice materials.

Team pricing, refunds and access

Enroll 2-20 participants together for 20% off. A full refund is available before the course or within 7 days after your selected cohort begins. Lifetime access applies to the course/community service; future editions require registration with Ion. Full terms are available before payment.

Earn your completion certificate

Submit at least four assignments, including the final improvement report, for instructor review. This is a completion certificate, not an accredited qualification.

Enroll or ask a question

smartaitomation.com/ai-evals-course
imosnoi@smartaitomation.com